

A short introduction to

Allmon

**a generic performance and
availability monitoring system**

Tomasz Sikora – London 2009

List of topics

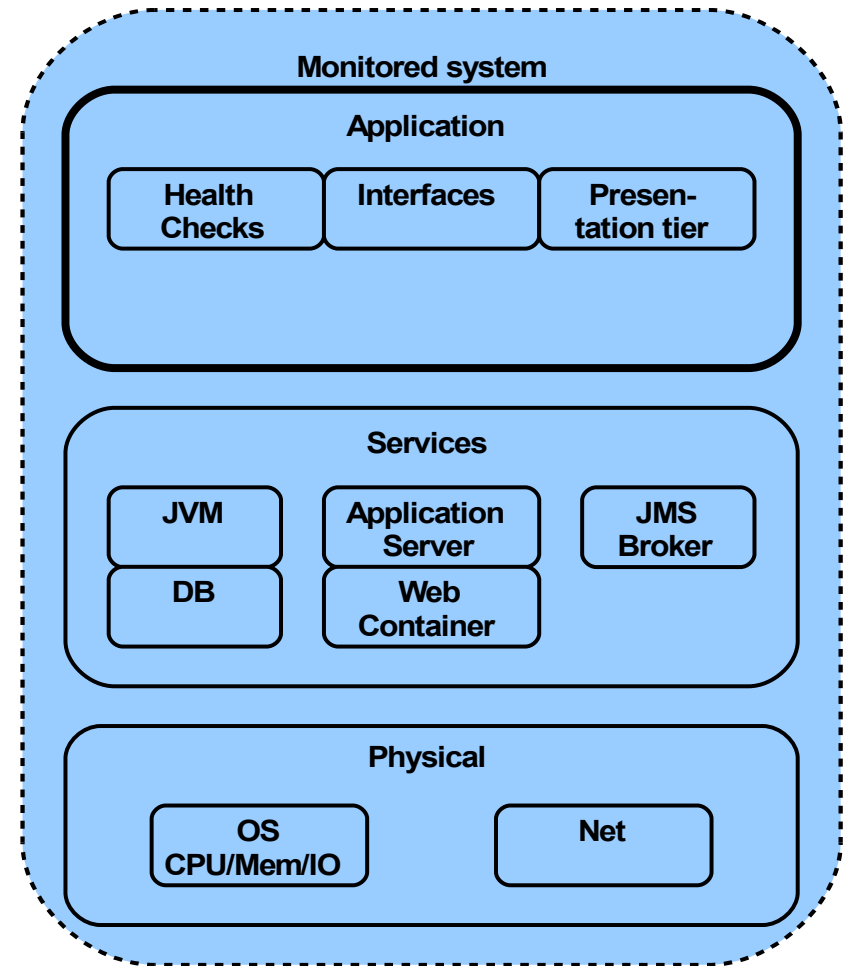
- Source of performance problems
- Continuous monitoring and metrics acquisition
- Allmon architecture (scalability, messaging)
- Deployment (distributed system)
- Configuration
- Analysis (use cases)
- Questions

Source of performance problems

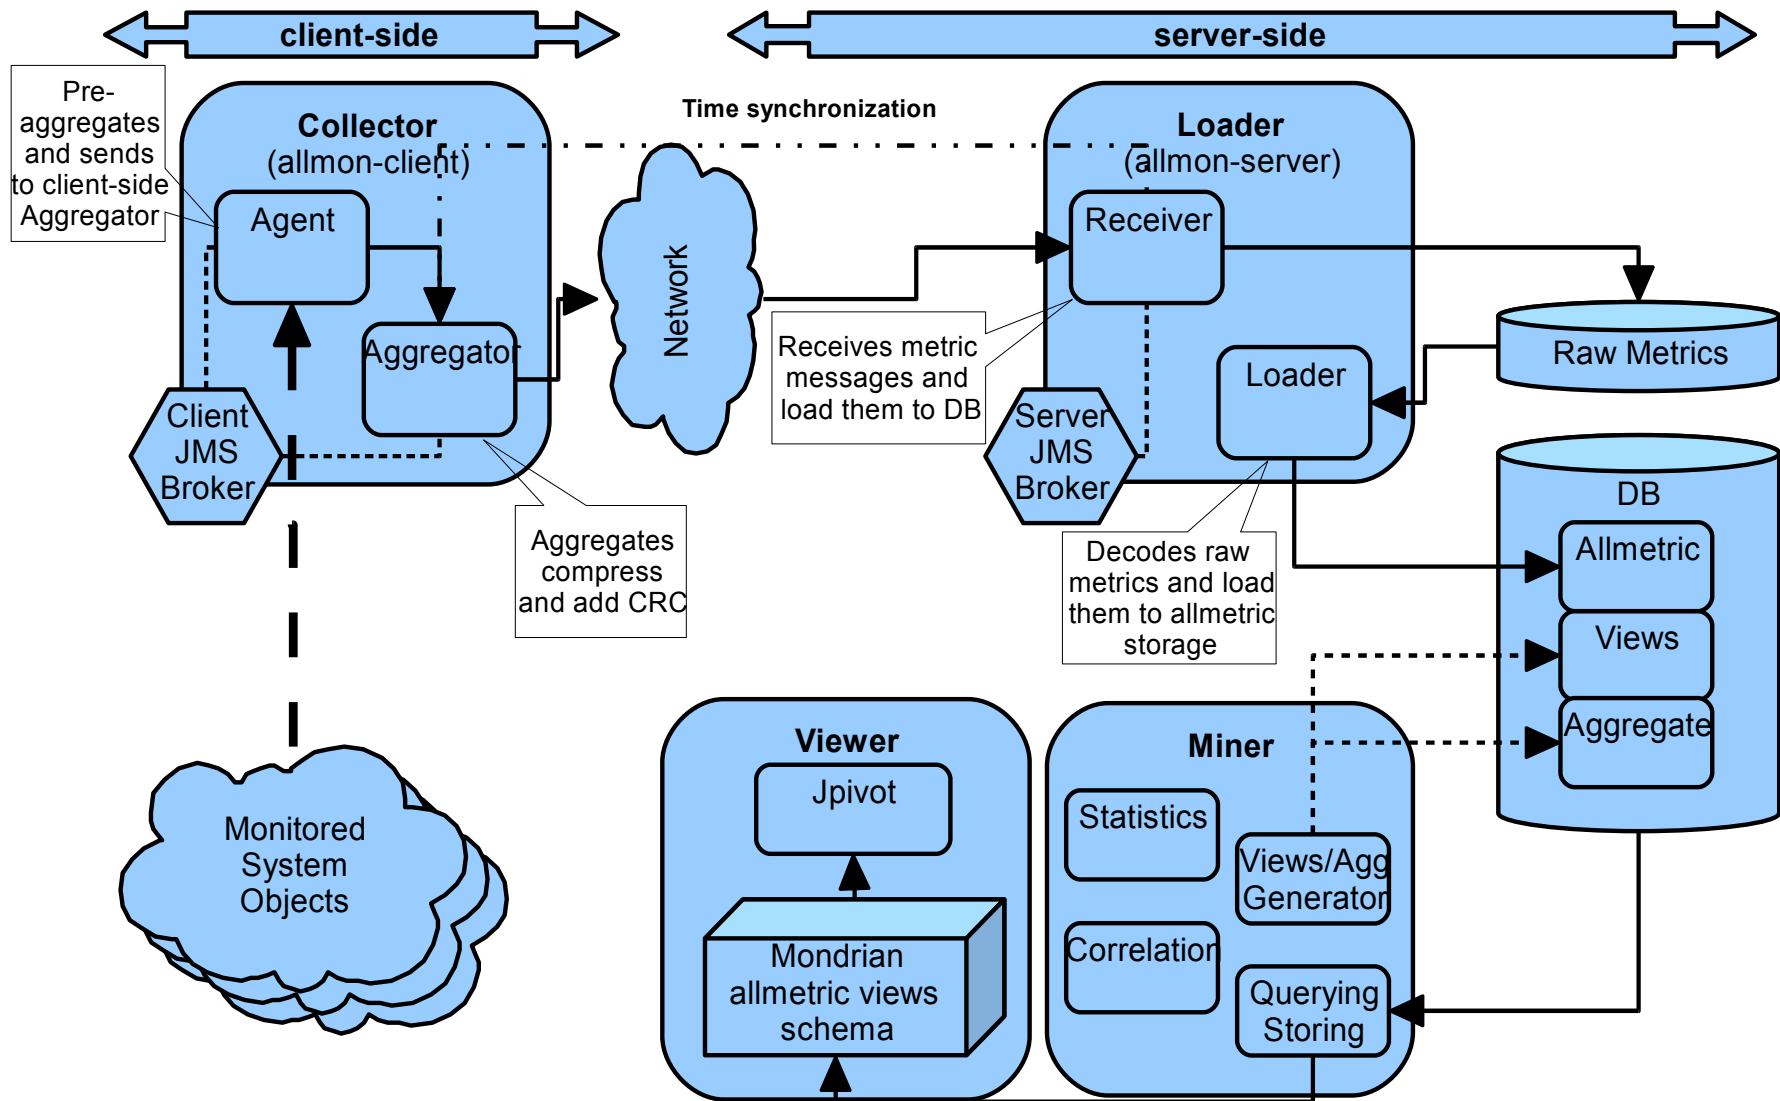
- **Lack of understanding (knowledge)**
 - Not well defined requirements (functional/non-functional)
 - Not experienced developers
 - Not well trained users
- **Lack of visibility**
 - Developers have to understand operational aspects
 - Multi-layer system monitoring is essential
- **Not enough testing**
 - Load tests (integration, regression) + Soak tests
 - Monitoring for "before-after" load test comparison and store all results

Continuous monitoring (metrics acquisition)

- **Monitoring multi-layer enterprise systems**
 - Main layers: Application, Services, Physical
 - Service Health Checks
- **Distributed system**
 - Messaging: isolation, non-intrusive, reliable
- **What to monitor, what to analyse?**
 - Collected data can be used for correlation analysis



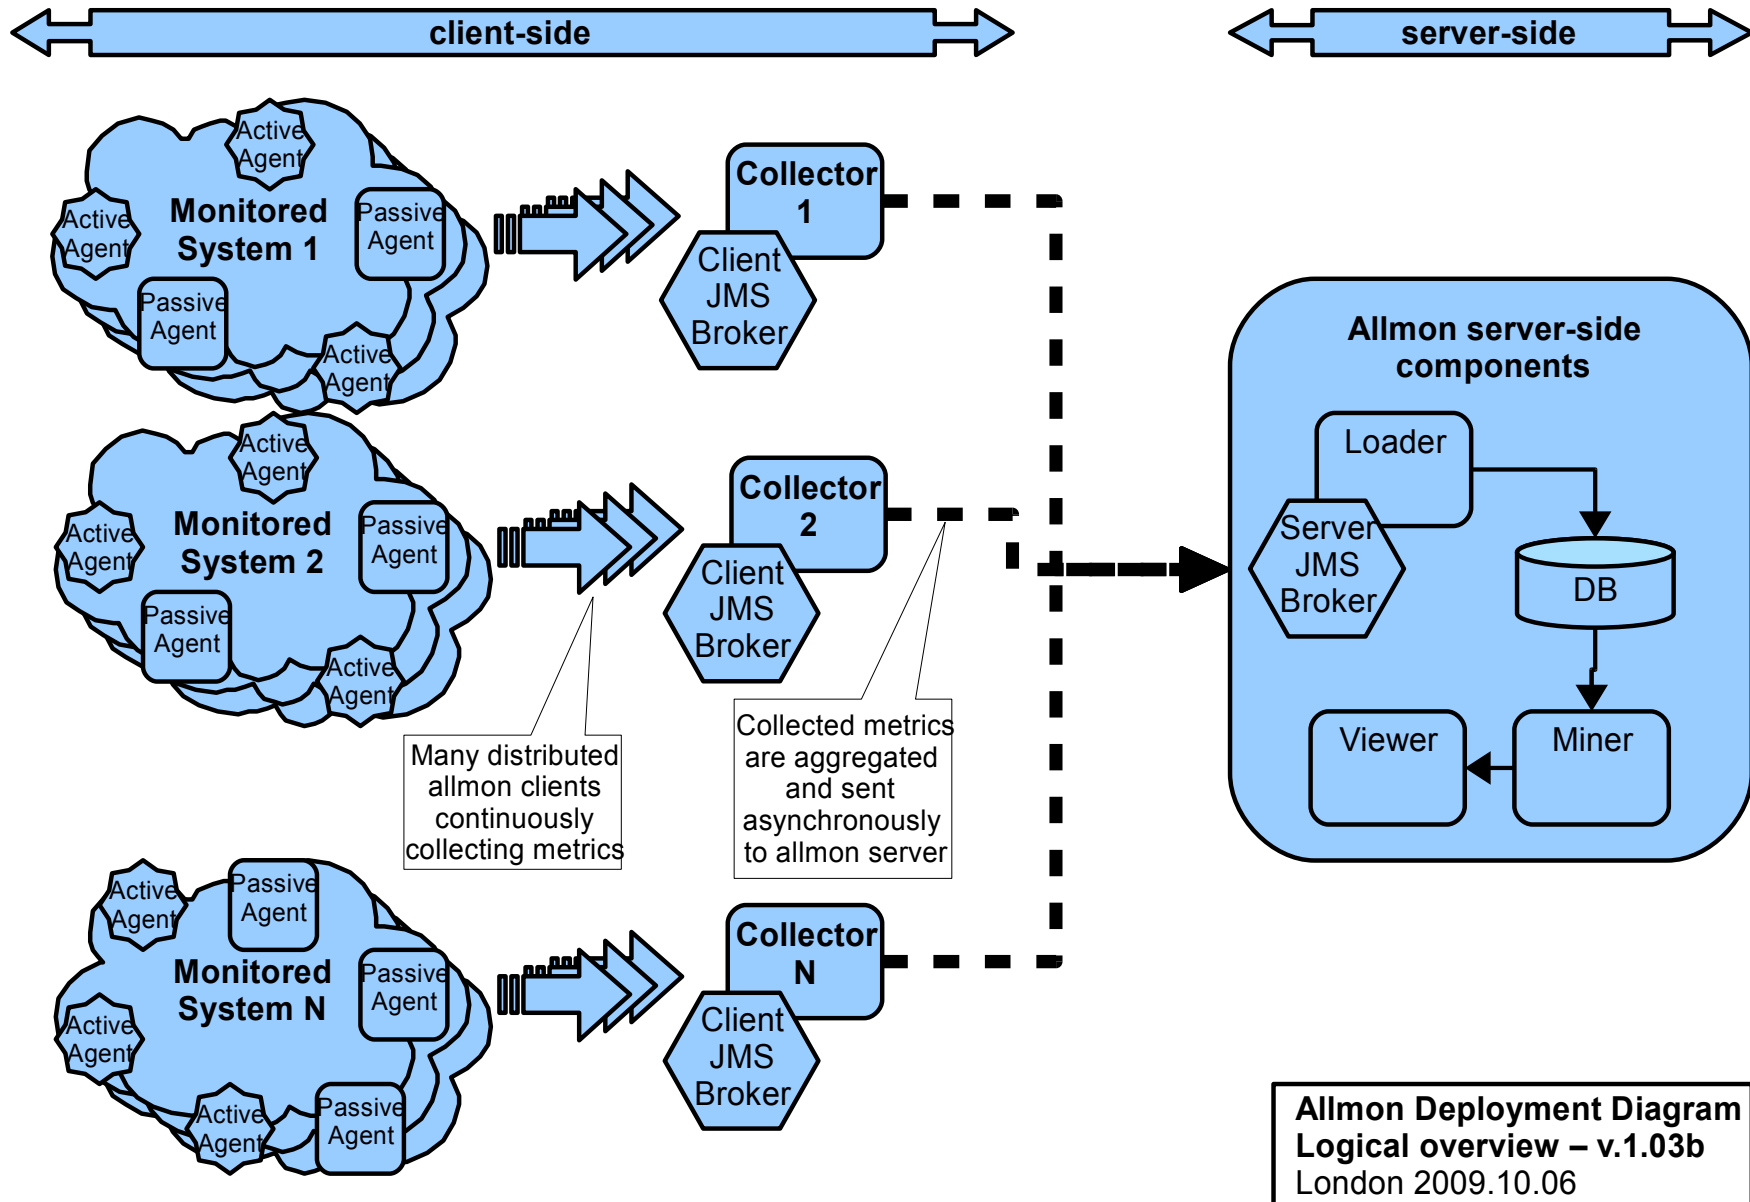
Allmon architecture



Allmon architecture

- **Collector** (distributed client-side)
 - Agents (Passive/Active agents collecting metrics)
 - Aggregator (common pre-aggregating and sending data mechanism)
- **Loader** (centralized server-side)
 - Miner (transforming data to allmetrics, aggregating)
 - Viewer (presentation, multidimensional analysis)
- **Data storage**
 - Raw data storage (staging tier)
 - Allmetric schema (generic 3NF structure)
 - Aggregates, pre-calculated structures (access tier)

Deployment



Configuration

- **Client-side**

- Agents configuration
 - Independent configuration for different types of agents
- Active agents scheduling
 - Based on crontab (cron4j)

- **Server-side**

- Database connectivity
- Loading process scheduling
- Aggregate processes parametrisation
- Visualization set-up

Allmon analysis - use cases

Collecting multi-tier system metrics is crucial for understanding system and finally finding performance problems (two simple examples)

- **Study case 1** - Growing JVM memory allocation (leaks) in comparison between several releases and users activity
 - Input: JVM metrics via JMX (mem, GC), application actions stats
 - Output: differences in users activity, differences in application behaviour and allocated resources
 - Action: identifying areas in code base responsible for huge deltas in memory consumption
- **Study case 2** - Not efficient interactions with services and databases
 - Input: database metrics, DB OS stats, application actions stats, intercepted persistence level calls
 - Output: rankings of: the longest performing application actions, the biggest product of execution count and execution times
 - Action: Easier prioritisation of areas which have to be improved

Questions

???

Allmon 2009

Allmon project page:
<http://code.google.com/p/allmon/>

Allmon user group:
<http://groups.google.com/group/allmon>

Code license: Apache License 2.0:
<http://www.apache.org/licenses/LICENSE-2.0>